

Lecture 7 Instance-based Learning

- Introduction of Instance-based learning
- K-Nearest Neighbor

Instance-based Learning - Overview

- In contrast to learning methods that construct a general, explicit description of the target function when training examples are provided, instance-based learning methods simply store the training examples.
- Generalizing beyond these examples is postponed until a new instance must be classified.
- Each time a new query instance is encountered, its relationship to the previously stored examples is examined in order to assign a target function value for the new instance.

Instance-based Learning - Overview

1. Instance-based learning methods are sometimes referred to as **delayed/lazy learning** methods because they delay processing until a new instance must be classified.
2. A key advantage of delayed/lazy learning is that instead of estimating the target function one for the entire instance space, these methods can estimate it **locally** and differently for each new instance to be classified.

Instance-based Learning - Overview

- Instance-based learning methods such as **Nearest Neighbor** are conceptually straightforward approaches to approximating real-valued or discrete-valued target functions.
 1. Learning in these algorithms consists of simply storing the presented training data.
 2. When a new query instance is encountered, a set of similar related instances is retrieved from memory and used to classify the new query instance.

Instance-based Learning - Overview

- Many techniques construct only a local approximation to the target function that applies in the neighborhood of the new query instance, and **never** construct an approximation designed to perform well over the entire instance space.
- This has significant advantages when the target function is very complex, but can still be described by a collection of less complex local approximations.

Instance-based Learning - Disadvantages

1. One disadvantage of instance-based approaches is that the cost of classifying new instances can be high.
 - This is due to the fact that nearly all computation takes place at classification time rather than when the training examples are first encountered.

Instance-based Learning - Disadvantages

2. A second disadvantage to many instance-based approaches, especially nearest neighbor approaches, is that they typically consider **all** attributes of the instances when attempting to retrieve similar training examples from memory.
 - If the target concept depends on only a few of the many available attributes, then the instances that are truly most "similar" may well be a large distance apart.

K-nearest Neighbor Learning (Intro.)

- The most basic instance-based method is the k-NEAREST NEIGHBOR algorithm.
- This algorithm assumes all instances correspond to points in the n-dimensional space \mathcal{R}^n . That means inputs of data are numeric ones.
- The nearest neighbors of an instance are defined in terms of the standard Euclidean distance.

K-nearest Neighbor Learning (Euclidean distance)

- Let an arbitrary instance \mathbf{x} be described by the feature vector $\langle a_1(\mathbf{x}), a_2(\mathbf{x}), \dots, a_n(\mathbf{x}) \rangle$
- where $a_r(\mathbf{x})$ denotes the value of the r^{th} attribute of instance \mathbf{x} . Then the distance between two instances \mathbf{x}_i and \mathbf{x}_j is defined to be $d(\mathbf{x}_i, \mathbf{x}_j)$, where

$$d(\mathbf{x}_i, \mathbf{x}_j) \equiv \sqrt{\sum_{r=1}^n (a_r(\mathbf{x}_i) - a_r(\mathbf{x}_j))^2}$$

K-nearest Neighbor Learning (output type)

- In nearest-neighbor learning the target function may be either discrete-valued or real-valued.
- Let us first consider learning discrete-valued target functions of the form $f : \mathcal{R}^n \rightarrow V$, where V is the finite set $\{v_1, \dots, v_s\}$.
- The k-NEAREST NEIGHBOR algorithm for approximating a discrete-valued target function is given in next page ->

K-nearest Neighbor Algorithm for approximating a discrete-valued function $f : \Re^n \rightarrow V$,

- Training algorithm:
 - For each training example $\langle x, f(x) \rangle$, add the example to the listing `training_examples`
- Classification algorithm
 - Given a query instance x_q to be classified,
 - Let x_1, \dots, x_k denote the k instances from `training_examples` that are nearest to x_q
 - Return



$$\hat{f}(x_q) \leftarrow \arg \max_{v \in V} \sum_{i=1}^k \delta(v, f(x_i))$$

- Where $\delta(a, b) = 1$ if $a = b$ and 0 otherwise.

K-nearest Neighbor Algorithm for approximating a discrete-valued function $f : \mathcal{R}^n \rightarrow V$,

- As shown there, the value $\hat{f}(x_q)$ returned by this algorithm as its estimate of $f(x_q)$ is just the most common value of f among the k training examples nearest to x_q .
- If we choose $k = 1$, then the 1-NEAREST NEIGHBOR algorithm assigns to $\hat{f}(x_q)$ the value $f(x_i)$ where x_i is the training instance nearest to x_q .
- For larger values of k , the algorithm assigns the most common value among the k nearest training examples.

K-nearest Neighbor Algorithm for approximating a discrete-valued function $f : \Re^n \rightarrow V$



1 Nearest Neighbor classifies x_q positive, 5 Nearest Neighbor classifies x_q as negative. The diagram on the right hand side shows the decision surface induced by 1-Nearest Neighbor.

K-nearest Neighbor Algorithm for approximating a discrete-valued function $f : \Re^n \rightarrow V$

- Calculate the mean value of the k nearest training examples.
- Replace the final line of previous algorithm by:

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k f(x_i)}{k}$$

Distance-Weighted Nearest Neighbor Algorithm for discrete-valued target functions

- Weight the contribution of each of the k neighbors according to their distance to the query point x_q , giving greater weight to closer neighbors.
- If x_q exactly match one of the training instances x_i , the denominator $d(x_q, x_i)^2$ is zero and, we assign $\hat{f}(x_q)$ to be $f(x_i)$.

$$\hat{f}(x_q) \leftarrow \arg \max_{v \in V} \sum_{i=1}^k w_i \delta(v, f(x_i))$$

where

$$w_i \equiv \frac{1}{d(x_q, x_i)^2}$$

Distance-Weighted Nearest Neighbor Algorithm for continuous- valued target functions

- We can distance-weight the instances for real-valued target functions in a similar way.

$$\hat{f}(x_q) \leftarrow \frac{\sum_{i=1}^k w_i f(x_i)}{\sum_{i=1}^k w_i}$$

where

$$w_i \equiv \frac{1}{d(x_q, x_i)^2}$$

Remark

- Note all of the above variants of the k-NEAREST NEIGHBOR algorithm consider only the k nearest neighbors to classify the query point.
- Once we add distance weighting, there is no harm in allowing all training examples because very distant examples will have very little effect on $\hat{f}(x_q)$
- The only disadvantage of considering all examples is that our classifier will run more slowly.
- If all training examples are considered when classifying a new query instance, we call the algorithm a global method.
- If only the nearest training examples are considered, we call it a local method.

Remark

- The distance-weighted k-NEAREST NEIGHBOR algorithm is a highly effective inductive inference method for many practical problems.
- It is robust to noisy training data and quite effective when it is provided a sufficiently large set of training data.
- Note that by taking the weighted average of the k neighbors nearest to the query point, it can smooth out the impact of isolated noisy training examples.

Remark

- Since the algorithm delays all processing until a new query is received, significant computation can be required to process each new query.
- Various methods have been developed for indexing the stored training examples so that the nearest neighbors can be identified more efficiently at some additional cost in memory.
- One such indexing method is the kd-tree
- https://www.ri.cmu.edu/pub_files/pub1/moore_andrew_1991_1/moore_andrew_1991_1.pdf